



DETR: End-to-End Object Detection with Transformers


By Sirak Ghebremusse
w251 Summer 2020

Topic

- Breakdown of research paper
- Training of DETR model with mixed precision
- Provide Docker files for cloud training and Jetson inference

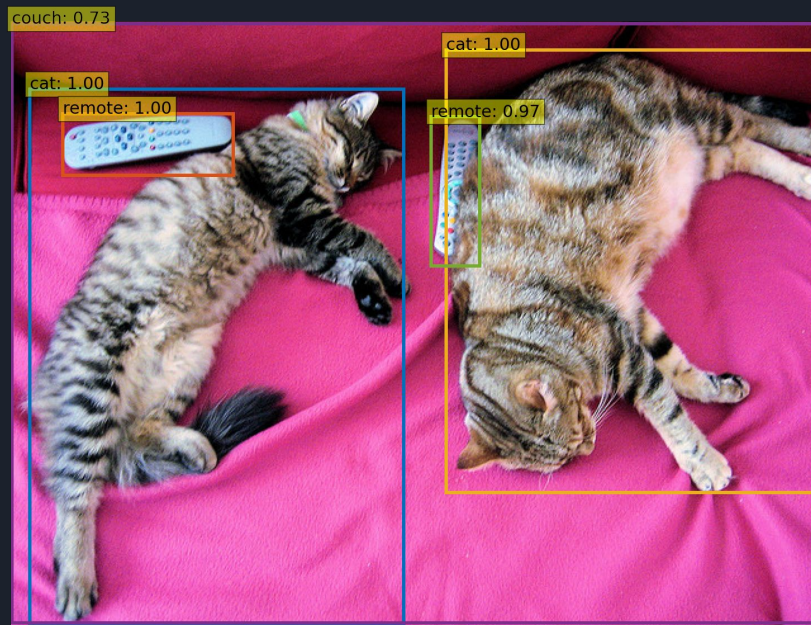
DE:TR: End-to-End Object Detection with Transformers

PyTorch training code and pretrained models for DETR (DEtection TRansformer). We replace the full complex hand-crafted object detection pipeline with a Transformer, and match Faster R-CNN with a ResNet-50, obtaining **42 AP** on COCO using half the computation power (FLOPs) and the same number of parameters. Inference in 50 lines of PyTorch.





Dataset

- Train and evaluation using the COCO dataset






Tools


- Pytorch
- TensorRT
- OpenCV
- Docker
- Nvidia Apex
- Jetson Xavier NX
- IBM and AWS Cloud VMs




DETR vs Yolo



Input image




Feature extraction network




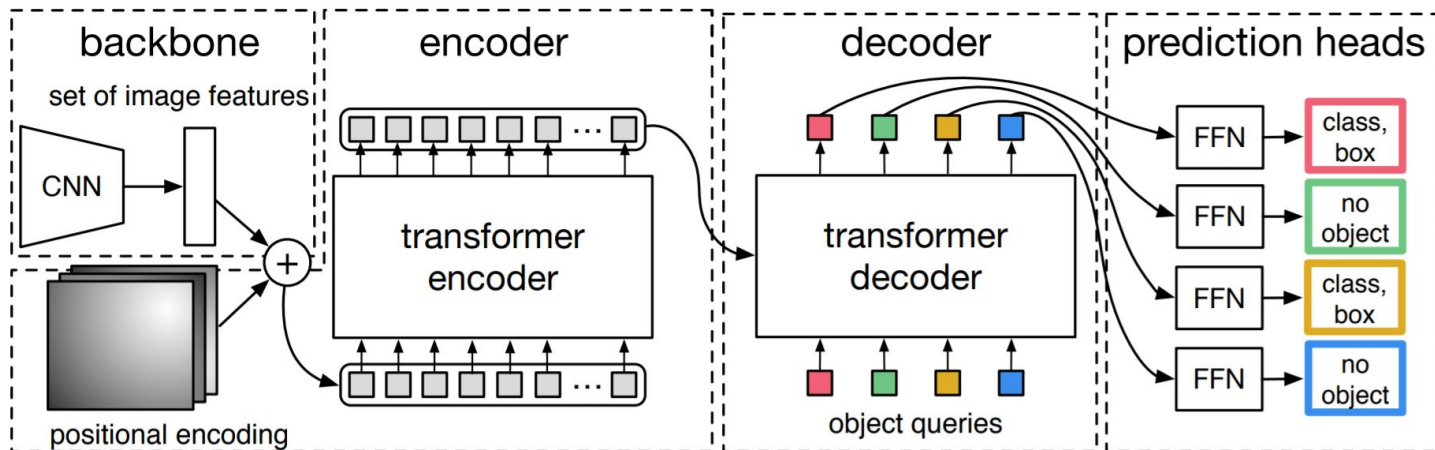
Bounding box prediction

NMS



Final detection

DETR vs Yolo



Attention for Detection





Fig. 7: Visualization of all box predictions on all images from COCO 2017 val set for 20 out of total $N = 100$ prediction slots in DETR decoder. Each box prediction is represented as a point with the coordinates of its center in the 1-by-1 square normalized by each image size. The points are color-coded so that green color corresponds to small boxes, red to large horizontal boxes and blue to large vertical boxes. We observe that each slot learns to specialize on certain areas and box sizes with several operating modes. We note that almost all slots have a mode of predicting large image-wide boxes that are common in COCO dataset.

Attention for Segmentation


self-attention(200, 200)




self-attention(280, 400)



self-attention(200, 600)



self-attention(440, 800)



Attention for Segmentation

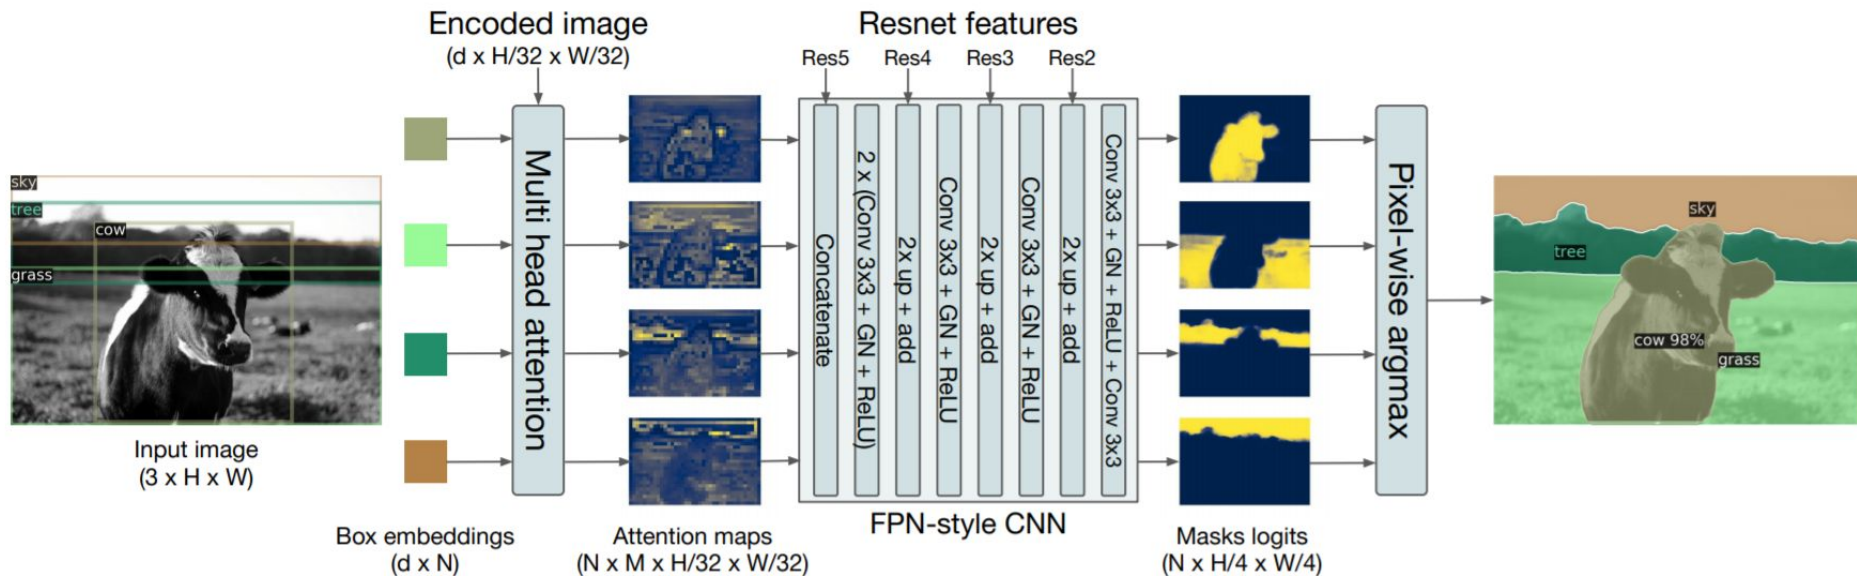


Fig. 8: Illustration of the panoptic head. A binary mask is generated in parallel for each detected object, then the masks are merged using pixel-wise argmax.



Mixed Precision Training

- 2 Docker Containers :
 - Download Coco (18 GB)
 - Train DETR model (install Apex)
- Training Modifications
 - Include Apex commands in scripts
 - Changed backbone to ResNet-18
 - Optimization level set to O1, for 16 bit weights



ubuntu@ip-172-31-42-183:~/detr\$ nvidia-smi

Sun Jul 26 17:46:39 2020



NVIDIA-SMI 440.33.01 Driver Version: 440.33.01 CUDA Version: 10.2									
GPU	Name	Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr.	ECC	
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute	M.	
0	Tesla V100-SXM2...	On	00000000:00:17:0	Off		0			
N/A	66C	P0	200W / 300W	15809MiB / 16160MiB		99%	Default		
1	Tesla V100-SXM2...	On	00000000:00:18:0	Off		0			
N/A	55C	P0	249W / 300W	15715MiB / 16160MiB		99%	Default		
2	Tesla V100-SXM2...	On	00000000:00:19:0	Off		0			
N/A	56C	P0	291W / 300W	15923MiB / 16160MiB		70%	Default		
3	Tesla V100-SXM2...	On	00000000:00:1A:0	Off		0			
N/A	65C	P0	191W / 300W	15949MiB / 16160MiB		61%	Default		
4	Tesla V100-SXM2...	On	00000000:00:1B:0	Off		0			
N/A	64C	P0	91W / 300W	15551MiB / 16160MiB		60%	Default		
5	Tesla V100-SXM2...	On	00000000:00:1C:0	Off		0			
N/A	56C	P0	78W / 300W	15455MiB / 16160MiB		60%	Default		
6	Tesla V100-SXM2...	On	00000000:00:1D:0	Off		0			
N/A	56C	P0	80W / 300W	15855MiB / 16160MiB		67%	Default		
7	Tesla V100-SXM2...	On	00000000:00:1E:0	Off		0			
N/A	67C	P0	87W / 300W	16027MiB / 16160MiB		75%	Default		

Processes:					GPU Memory
GPU	PID	Type	Process name		Usage
0	2897	C	.../anaconda3/envs/pytorch_p36/bin/python3		15797MiB
1	2898	C	.../anaconda3/envs/pytorch_p36/bin/python3		15703MiB
2	2899	C	.../anaconda3/envs/pytorch_p36/bin/python3		15911MiB
3	2900	C	.../anaconda3/envs/pytorch_p36/bin/python3		15937MiB
4	2901	C	.../anaconda3/envs/pytorch_p36/bin/python3		15539MiB
5	2902	C	.../anaconda3/envs/pytorch_p36/bin/python3		15443MiB
6	2903	C	.../anaconda3/envs/pytorch_p36/bin/python3		15843MiB
7	2904	C	.../anaconda3/envs/pytorch_p36/bin/python3		16015MiB



ubuntu@ip-172-31-42-183:~/detr\$

Accumulating evaluation results...

DONE (t=10.07s).

IoU metric: bbox

Average Precision (AP) @[IoU=0.50:0.95	area=	all	maxDets=100	= 0.338
Average Precision (AP) @[IoU=0.50	area=	all	maxDets=100	= 0.546
Average Precision (AP) @[IoU=0.75	area=	all	maxDets=100	= 0.345
Average Precision (AP) @[IoU=0.50:0.95	area=	small	maxDets=100	= 0.135
Average Precision (AP) @[IoU=0.50:0.95	area=	medium	maxDets=100	= 0.368
Average Precision (AP) @[IoU=0.50:0.95	area=	large	maxDets=100	= 0.527
Average Recall (AR) @[IoU=0.50:0.95	area=	all	maxDets= 1	= 0.290
Average Recall (AR) @[IoU=0.50:0.95	area=	all	maxDets= 10	= 0.461
Average Recall (AR) @[IoU=0.50:0.95	area=	all	maxDets=100	= 0.500
Average Recall (AR) @[IoU=0.50:0.95	area=	small	maxDets=100	= 0.224
Average Recall (AR) @[IoU=0.50:0.95	area=	medium	maxDets=100	= 0.546
Average Recall (AR) @[IoU=0.50:0.95	area=	large	maxDets=100	= 0.758

Training time 2 days, 1:56:26

real 2997m8.389s
user 26195m47.450s
sys 6497m46.220s

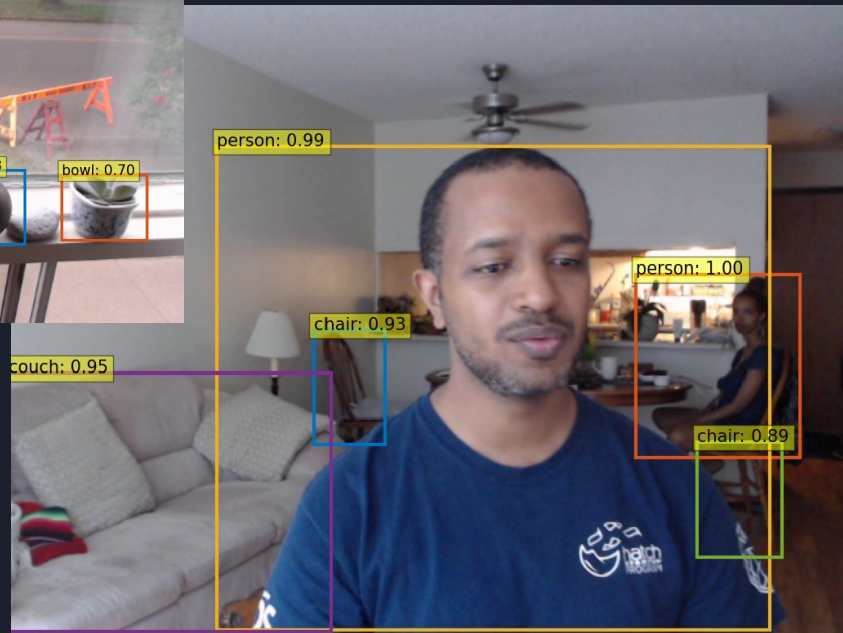


(pytorch_p36) ubuntu@ip-172-31-42-183:~/detr\$ q





Mixed Precision Inference

- Torch2TensorRT container
- Inference container for TensorRT or Pytorch model

```
205
206         # propagate through the model
207         sample = transform(img).unsqueeze(0)
208         #sample.cuda()
209         sample.to(device)
210         outputs = model(sample)
211
212         # keep only predictions with 0.7+ confidence
213         probas = outputs['pred_logits'].softmax(-1)[0, :, :-1]
214         keep = probas.max(-1).values > 0.7
215
216         # convert boxes from [0; 1] to image scales
217         bboxes_scaled = rescale_bboxes(outputs['pred_boxes'][0, keep], img.size)
218
219         # Display the resulting frame
220         plot_results(img, probas[keep], bboxes_scaled)
221
```



bus: 0.95





Challenges

- Torch version and GPU driver mismatch on IBM
- APEX installation required NGC containers
- Unable to fully utilize AWS GPUs (~70%)
- Torch2TensorRT not very intuitive
-



Future Goals

- Train for Image segmentation
- Multi-server training
- Cast layers to below 16-bit precision
- Ensure it runs on Tensor cores on Jetson NX
- Smaller vision backbone!



Thank you!



Links

<https://ai.facebook.com/blog/end-to-end-object-detection-with-transformers/>

<https://ai.facebook.com/research/publications/end-to-end-object-detection-with-transformers/>

https://www.youtube.com/watch?v=T35ba_VXkMY&t=1464s

https://colab.research.google.com/github/facebookresearch/detr/blob/colab/notebooks/detr_demo.ipynb

https://colab.research.google.com/github/facebookresearch/detr/blob/colab/notebooks/detr_demo.ipynb

https://colab.research.google.com/github/facebookresearch/detr/blob/colab/notebooks/DETR_panoptic.ipynb

Github Repo: <https://github.com/sirakzg/detr>